

Programovanie v operačných systémoch

04 - Memory

Jozef Šiška



Department of Applied Informatics
Comenius University in Bratislava

2020/2021

Memory management

Memory management

- ▶ Allocation
 - ▶ kernel: brk, mmap
 - ▶ C/C++: malloc, realloc, free, mmap, new, delete
- ▶ "Management"
 - ▶ pairing alloc/release, memory leaks
 - ▶ ownership, passing between functions etc. (size?)
 - ▶ dangling pointers
 - ▶ `std::unique_ptr`, `std::shared_ptr`
- ▶ Reference counting
 - ▶ RAIL, immediate release, cycles?
 - ▶ implicit sharing, COW
- ▶ Garbage collection
 - ▶ when will it happen? price of detection?

Reference counting

- ▶ `std::shared_ptr`
- ▶ immediate "release", RAII similar to other resources
- ▶ cheap / fast (at least relatively: large object "trees" can take a while to release, which can be noticable in realtime apps)
- ▶ slight space (refcount/control block) / speed (inc/dec) overhead
- ▶ memory access - one or two dereferences
- ▶ synchronization (atomic refcount)
- ▶ cycles!
- ▶ breaking cycles: weak references
 - ▶ can become dangling
 - ▶ reference "zeroing"
 - ▶ keep track of both weak and strong references
 - ▶ when "strong" refcount becomes zero, data is released and weak references can't be used anymore to access data
 - ▶ `std::shared_ptr + std::weak_ptr`

- ▶ doesn't combine nicely with management of other resources... (Java `finalize()`)
- ▶ "unpredictable", performance...
- ▶ reference counting + cycle detection (Python)
- ▶ tracing - find objects not reachable from "root" objects

So... how to avoid memory leaks in C / bad C++ / ...?

- ▶ release resources before each **return**
- ▶ **goto** solutions
- ▶ with exceptions in C++, everything is a possible **return**!
- ▶ valgrind (memcheck)
- ▶ (and other tools...)

... and is Java really safe?

- ▶ "hidden" references: registering listeners, observers,...

Copy on write (COW)

- ▶ reference counting on steroids
 - ▶ cheap pass by value even for very large objects
 - ▶ don't make copies when not needed
 - ▶ shared data - every data class is basically a shared (refcounted) pointer
 - ▶ Copy on change
 - ▶ C++: const vs non-const methods
 - ▶ when refcount > 1
 - ▶ might not be always possible/feasible (std::string in C++11?)
 - ▶ unpredictable/unintuitive complexity/efficiency
- ```
String s2 = s1; s2[0] = 'a';
```

# Small string/object optimisation

- ▶ Container classes allocate data buffers on heap
- ▶ Usually need data pointer, capacity and size  
`sizeof(void*) + sizeof(size_t) + sizeof(size_t) == 24;`
- ▶ If the data is small, this is a "waste" (24 bytes on stack to point to 2 byte string on heap)
- ▶ Sacrify one bit from size/capacity (or maybe pointer) as a flag that the fields contains actual data