

Programovanie v operačných systémoch

05 - Network

Jozef Šiška



Department of Applied Informatics
Comenius University in Bratislava

2020/2021

Networking Recap

- ▶ HW
- ▶ Ethernet
- ▶ IP
 - ▶ TCP
 - ▶ UDP
 - ▶ DNS
- ▶ Local sockets (Unix sockets)

Socket

Create a socket (not yet connected to anything)

```
int socket(int domain, int type, int protocol);
```

Domain

AF_UNIX, AF_LOCAL	local sockets (man 7 unix)
AF_INET	IPv4 (man 7 ip)
AF_INET6	IPv6 (man 7 ipv6)
AF_NETLINK	kernel userspace interface (man 7 netlink)
AF_IPX, AF_X25, ...	

Type

SOCK_STREAM	reliable byte stream (i.e. TCP)
SOCK_DGRAM	connectionless, unreliable messages (UDP)
SOCK_SEQPACKET, SOCK_RAW, ...	

man 2 socket

man 7 {unix,ip,ipv6}

Connect and communicate

Server

- ▶ Bind a socket to an address

```
int bind(int sockfd, const struct sockaddr *addr,  
         socklen_t addrlen);
```

- ▶ Listen on the socket for incoming connections

```
int listen(int sockfd, int backlog);
```

- ▶ Accept a connection

```
int accept(int sockfd, struct sockaddr *addr,  
           socklen_t *addrlen);
```

Client

```
int connect(int sockfd, const struct sockaddr *addr,  
            socklen_t addrlen);
```

man 2 {bind,listen,accept,connect}

Reading, writing, closing

Read	Write	
read	write	plain read/write
recv	send	specify additional flags
recvfrom	sendto	get / specify peer address (i.e. UDP packets)
recvmsg	sendmsg	readv/writev style, additional data
shutdown		close (one direction of) a connection
close		close (dispose of) the socket

Addresses

A general "some address" type (man 2 bind):

```
struct sockaddr {
    sa_family_t sa_family;
    char        sa_data[14];
}
```

IPv4 address (man 7 ip, IPv6 is similar):

```
struct sockaddr_in {
    sa_family_t    sin_family; /* address family: AF_INET */
    in_port_t      sin_port;   /* port in network byte order */
    struct in_addr sin_addr;   /* internet address */
};

/* Internet address. */
struct in_addr {
    uint32_t s_addr; /* address in network byte order */
};
```

Need to cast between types:

```
struct sockaddr_in  addr;
/* set the fields, open socket */
ret = bind(sockfd, (struct sockaddr *) &addr, sizeof(addr));
```

Obtaining, printing addresses

Any address (for server)

```
struct sockaddr_in addr;  
addr.sin_family = AF_INET;  
addr.sin_addr.s_addr = INADDR_ANY;
```

Network vs host order (ports)

```
uint16_t portno = 1234;  
addr.sin_port = htons(portno);  
portno = ntohs(addr.sin_port);
```

Convert IPv4 address to sockaddr_in

```
ret = inet_aton("127.0.0.1", &addr.sin_addr);  
ret = inet_pton(AF_INET, "127.0.0.1", &addr.sin_addr);  
ret = inet_pton(AF_INET6, "127.0.0.1", &addr.sin6_addr);
```

Convert sockaddr_in to IPv4 address

```
printf("%s\n", inet_ntoa(addr.sin_addr));  
char str[INET_ADDRSTRLEN];  
ret = inet_ntop(AF_INET, &addr.sin_addr, str, len);
```


Resolving DNS addresses

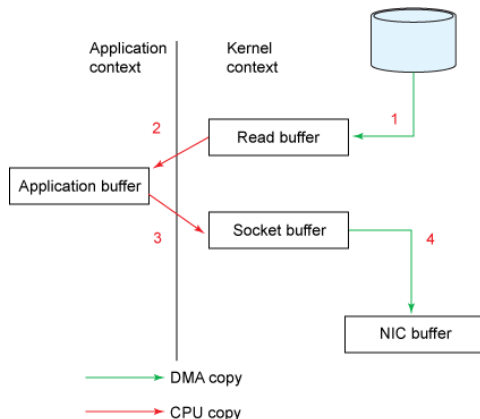
gethostbyname

```
struct sockaddr_in addr; struct hostent *server;
server = gethostbyname(str);
if (!server) { /* ... */ };
memcpy(server->h_addr, &addr.sin_addr.s_addr, server->h_length);
```

getaddrinfo

```
struct addrinfo hints, *result, *rp;
int sockFd = -1;
int ret = getaddrinfo("www.fmph.uniba.sk", "http", &hints, &result);
if (ret != 0) { /*...*/ }
for (rp = result; rp != null; rp = rp->ai_next) {
    sockFd = socket(rp->ai_family, rp->ai_socktype, rp->ai_protocol));
    if (sockFd == -1)
        continue;
    if (connect(sockFd, rp->ai_addr, rp->ai_addrlen) == 0)
        break;
    close(sockFd);
}
if (rp == NULL) { /* could not connect to any of the addresses*/ }
freeaddrinfo(result);
// connected...
```

Copying - problems



[https:](https://www.ibm.com/developerworks/library/j-zero-copy/index.html)

[//www.ibm.com/developerworks/library/j-zero-copy/index.html](https://www.ibm.com/developerworks/library/j-zero-copy/index.html)

<https://www.linuxjournal.com/article/6345>

Copying - sendfile

```
ssize_t sendfile(int out_fd, int in_fd, off_t *offset, size_t count);
```

- ▶ in_fd must be "mmap-able"

Copying - sendfile

```
ssize_t sendfile(int out_fd, int in_fd, off_t *offset, size_t count);
```

- ▶ in_fd must be "mmap-able"

```
ssize_t splice(int fd_in, loff_t *off_in, int fd_out,  
               loff_t *off_out, size_t len, unsigned int flags);
```

- ▶ At least one of fd_in and fd_out must be a pipe.
- ▶ See also man 2 tee, man 2 vmsplice